

Project 1 in FMNN10 and NUMN20

© Tony Stillfjord, Monika Eisenmann, Gustaf Söderlind 2021

Goals. The goal is to get started with MATLAB or Python as a tool for approximating solutions to differential equations, and in particular, to get acquainted with time-stepping methods for initial value ODEs, and standard techniques for analyzing and assessing such methods.

Contents:

Explicit Runge–Kutta methods and automatic step size control. You construct your own ODE solver, based on an explicit Runge–Kutta method with embedded error estimator. The error estimator is used to adjust the time step h along the integration, so that the error estimate is kept close to a prescribed accuracy tolerance `tol`.

Stiff vs. nonstiff problems. We focus on understanding the distinction between stiff and nonstiff problems, and why stiff problems require *implicit* methods with unbounded stability regions. We will study two nonlinear oscillatory systems, the nonstiff Lotka-Volterra population dynamics problem and the van der Pol equation, which may be stiff or nonstiff depending on problem parameters. You will work both with your own solver, and a professionally implemented stiff ODE solver.

Presentation: You work individually or (preferably) with a fellow student on the programs. In week 3, you will present your code and results to one of the teaching assistants. Instructions for how to book a time slot for this will be posted on the course website. There is no fixed form for this presentation, but you need to present your work in a coherent way. We recommend that you create a few slides in, e.g., Beamer or Powerpoint with the main interesting code snippets and plots of the results. The plots need to have proper labels etc., but you should also explain how you interpret them and whether they show what you expected to see. The teaching assistant will ask a few follow-up questions. In most cases, any misunderstandings can be cleared up during the presentation, but if there are apparent large knowledge gaps you will need to redo the presentation one week later.

General advice: Prepare well before you go to the exercise. Read the entire instruction and work out a plan for how to solve the problems. *Get started as soon as possible.* The teaching assistants can only advise you when you have a program to work with. It is also highly recommended to do Project 0 first.

1. Explicit adaptive Runge–Kutta methods

Theory. An explicit Runge–Kutta method for the initial value problem $y' = f(t, y)$ is a method of the form exemplified by the classical 4th-order Runge–Kutta method (also known as RK4)

$$\begin{aligned} Y_1' &= f(t_n, y_n) \\ Y_2' &= f(t_n + h/2, y_n + hY_1'/2) \\ Y_3' &= f(t_n + h/2, y_n + hY_2'/2) \\ Y_4' &= f(t_n + h, y_n + hY_3') \\ y_{n+1} &= y_n + \frac{h}{6} (Y_1' + 2Y_2' + 2Y_3' + Y_4'). \end{aligned}$$

A single step of the method can then be described as follows. The method “samples” the right-hand side $f(t, y)$ at four different points to compute the four **stage derivatives** Y_1' , Y_2' , Y_3' and Y_4' . Then it forms a linear combination of these derivatives to obtain the “average derivative” to advance the solution from y_n to y_{n+1} .

A general Runge–Kutta method is defined by its coefficients a_{ij} for evaluating the Y_i' , and the coefficients b_j for forming the linear combination of these derivatives to update the solution. The method can be written

$$\begin{aligned} hY_i' &= hf(t_n + c_i h, y_n + \sum_{j=1}^s a_{ij} hY_j'), \quad i = 1, \dots, s, \\ y_{n+1} &= y_n + \sum_{j=1}^s b_j hY_j'. \end{aligned}$$

We see that the method is represented by two *coefficient vectors*, c and b , and the *coefficient matrix* A , usually arranged in the **Butcher tableau**

$$\begin{array}{c|c} c & A \\ \hline y & b^T \end{array}$$

For the classical RK4 method, the Butcher tableau is

$$\begin{array}{c|cccc} 0 & 0 & 0 & 0 & 0 \\ 1/2 & 1/2 & 0 & 0 & 0 \\ 1/2 & 0 & 1/2 & 0 & 0 \\ 1 & 0 & 0 & 1 & 0 \\ \hline y & 1/6 & 1/3 & 1/3 & 1/6 \end{array}$$

Most methods used in practice have $c_i = \sum_j a_{ij}$, so it is sufficient to know the matrix A and the vector b . Check that this condition holds for RK4.

Task 1.1 Write a MATLAB function

```
unew = RK4step(f, told, uold, h)
```

that takes a single step with the classical RK4 method. Here \mathbf{f} is the function defining the differential equation. Then test it by solving a simple problem of your choice. For example, you could use the linear test equation $y' = \lambda y$, and **verify that the global error is $O(h^4)$ by plotting the error in a log-log diagram**, like you did for the Euler methods in Project 0. This check ensures that your implementation is correct. Note that you now need to define the differential equation (\mathbf{f}) in a separate MATLAB m-file.

Theory. In a similar manner, a 3rd order RK method called RK3 is defined by the Butcher tableau

$$\begin{array}{c|ccc} 0 & 0 & 0 & 0 \\ 1/2 & 1/2 & 0 & 0 \\ 1 & -1 & 2 & 0 \\ \hline z & 1/6 & 2/3 & 1/6 \end{array}$$

Here we see that some of the evaluations of the right-hand side f are the same as for the classical RK4 method. In fact, we can write both methods *simultaneously* as

$$\begin{aligned} Y'_1 &= f(t_n, y_n) \\ Y'_2 &= f(t_n + h/2, y_n + hY'_1/2) \\ Y'_3 &= f(t_n + h/2, y_n + hY'_2/2) \\ Z'_3 &= f(t_n + h, y_n - hY'_1 + 2hY'_2) \\ Y'_4 &= f(t_n + h, y_n + hY'_3) \\ y_{n+1} &= y_n + \frac{h}{6} (Y'_1 + 2Y'_2 + 2Y'_3 + Y'_4) \\ z_{n+1} &= y_n + \frac{h}{6} (Y'_1 + 4Y'_2 + Z'_3). \end{aligned}$$

With *five* evaluations of the right-hand side f instead of four (the extra evaluation being Z'_3), we can obtain *both* a 3rd order approximation z_{n+1} and a 4th order approximation y_{n+1} from the same starting point y_n . This can be used to *estimate a local error by the difference* $\ell_{n+1} := z_{n+1} - y_{n+1}$.

When two methods use the same function evaluations, we say that we have an **embedded pair** of RK methods. The embedded pair above is

called RK34. In practice, one doesn't compute z_{n+1} . Instead, one computes the error estimate directly from

$$\ell_{n+1} := \frac{h}{6} (2Y'_2 + Z'_3 - 2Y'_3 - Y'_4).$$

This is not only to save one unnecessary computation, but because subtracting two very similar values is sensitive to round-off errors and can lead to an inaccurate estimate of the error.

Task 1.2 Starting from your `RK4step`, write a MATLAB function

`[unew, err] = RK34step(f, told, uold, h)`

that takes a single step with the classical RK4 method and puts the result in `unew`, and uses the embedded RK3 to compute *a local error estimate* in `err` as described above.

Theory. Let us use the Euclidean norm throughout and assume that the local error $r_{n+1} := \|\ell_{n+1}\|_2$ is of the form

$$r_{n+1} = \varphi_n h_n^k$$

if the step size h_n was used. The coefficient φ_n is called the **principal error function** and depends on the method as well as on the problem.

Our goal is now to keep $r_n = \text{TOL}$ for a prescribed accuracy tolerance TOL . The idea is to vary the step size so that **the magnitude of the local error remains constant along the solution**. In order to do this, let us assume that φ_n varies slowly (treat it as if it were “constant”).

Now suppose that r_n was a bit off, i.e., there is a deviation between r_n and the desired value TOL . How would we change the step size in order to eliminate this deviation? If φ is constant we seek a step size h_n such that

$$\begin{aligned} r_n &= \varphi h_{n-1}^k \\ \text{TOL} &= \varphi h_n^k, \end{aligned}$$

where the second equation says that the step size has been changed to h_n so that the error becomes equal to TOL in magnitude. As we know the old step size h_{n-1} , as well as the estimated error r_n and the tolerance TOL , we can solve for the next step size h_n . Thus, eliminating φ from the equations above by dividing the equations, we find

$$h_n = \left(\frac{\text{TOL}}{r_n} \right)^{1/k} \cdot h_{n-1}.$$

This is the simplest recursion for controlling the step size and is the desired algorithm for making the RK34 method **adaptive**. The power k is the order of the error estimator. Because RK4 has local error $O(h^5)$ and RK3 has local error $O(h^4)$, the difference between the two methods is $O(h^4)$. This means that you should use the value $k = 4$ for your adaptive RK34 method.

This simple step size controller is, however, not particularly good. Turning to control theory for better alternatives, you will use a proportional–integral controller (PI controller) with your RK34 code,

$$h_n = \left(\frac{\text{TOL}}{r_n} \right)^{2/(3k)} \left(\frac{\text{TOL}}{r_{n-1}} \right)^{-1/(3k)} \cdot h_{n-1}, \quad (1)$$

which is more robust. Here we use the current error estimate, as well as the previous error estimate. On the very first step, however, no “previous” error estimate is available. For $n = 1$ we then put $r_0 = \text{TOL}$, after which the recursion will start operating as intended.

Task 1.3 Write a function

```
hnew = newstep(tol, err, errold, hold, k)
```

which, given the tolerance `tol`, a local error estimate `err` and a previous error estimate `errold`, the old step size `hold`, and the order `k` of the error estimator, computes the new step size `hnew` using (1).

Task 1.4 Combining RK34 and `newstep`, write an adaptive ODE solver

```
[t,y] = adaptiveRK34(f, t0, tf, y0, tol)
```

which solves $y' = f(t, y)$; $y(t_0) = y_0$ on the interval $[t_0, t_f]$, while keeping the error estimate equal to `tol` using the step size control algorithm you implemented above. In the vector `t` you store the time points the method uses, and in `y` you store the corresponding numerical approximation, as a row vector for each value of `t`. See below how you need to arrange the function `f`. Make sure that in its last step, the method *exactly hits the end point* `tf`. Thus, in the last step your solver `adaptiveRK34` has to “override” the value `hnew` supplied by `newstep`.

In order to start the integration, you also need to pick an initial step size. We suggest using the formula

$$h_0 = \frac{|t_f - t_0| \cdot \text{TOL}^{1/4}}{100 \cdot (1 + \|f(y_0)\|)}.$$

You can start with a tolerance `tol` = 10^{-6} . What changes for different values of `tol`?

2. A nonstiff problem

Theory. A classical model in biological population dynamics is the **Lotka–Volterra equation**,

$$\begin{aligned}\dot{x} &= ax - bxy \\ \dot{y} &= cxy - dy,\end{aligned}$$

where a, b, c, d are positive parameters. The equation models the interaction between a predator species, y (foxes), and a prey, x (rabbits). If no foxes are present ($y = 0$) the rabbits multiply and grow exponentially. On the other hand, if there are no rabbits ($x = 0$), the foxes have no food supply and die at a rate determined by d . The product term, xy , represents the probability that a fox encounters a rabbit within their shared ecosystem. This encounter benefits the fox, which eats the rabbit, so the product term is positive in the second (fox) equation, and negative in the first (rabbit) equation.

The Lotka–Volterra equation is *separable*. By dividing the two equations, we get

$$\frac{dx}{dy} = \frac{ax - bxy}{cxy - dy} = \frac{x(a - by)}{y(cx - d)}.$$

Written in terms of differentials, we have

$$\left(c - \frac{d}{x}\right) dx = \left(\frac{a}{y} - b\right) dy,$$

and by integration we obtain $cx - d \log x = a \log y - by + K$. Hence the function

$$H(x, y) = cx + by - d \log x - a \log y$$

remains constant at all times, i.e., $H(x, y)$ is **invariant** along solutions. This means (non-trivially, proof omitted) that the Lotka–Volterra equation has **periodic solutions**.

Task 2.1 Choose the parameters a, b, c, d as $(3, \quad 9, \quad 15, \quad 15)$ and pick some suitable positive initial values, preferably not too far from the equilibrium point $(d/c, a/b)$, e.g. $(1, 1)$. Use your own adaptive RK34 solver to solve the problem. Run with a tolerance of (say) 10^{-6} or 10^{-8} . (You can use tighter tolerances still, if your code is good enough.) Simulate the system for at least 10 full periods.

Are the solutions periodic as claimed, and how long, approximately, is the period? Plot x and y as functions of time, and plot y as a function of x , i.e., the *phase portrait*. It is often easier to check periodicity there, because

the latter plot is a closed orbit if the solution is periodic. You may have to experiment a little with initial conditions and how long time you integrate in order to get a nice plot.

Investigate what happens if you change the initial conditions. Do you get the same periodic solution and does the period remain the same?

Integrate over a very long time (100 – 1,000 full periods, depending on choice of tolerance) to check whether the numerically computed $H(x, y)$ stays near its initial value $H(x(0), y(0))$ or drifts away. This can be done by plotting

$$|H(x, y)/H(x(0), y(0)) - 1|$$

as a function of time, where you insert the computed values of x and y into the expression. Choose a suitable type of plot. Should it be loglog or linlog?

Hints. In order to solve the Lotka–Volterra equation, you need to write a function that computes the right-hand side of the ODE in the following format,

```
function dudt = lotka(t,u)
a = ... ;
b = ... ;
c = ... ;
d = ... ;
dudt = [ a*u(1)-b*u(1)*u(2); c*u(1)*u(2)-d*u(2) ];
```

Note that the argument t is necessary, although it will not be used by your solver in this case. The reason is that this format is used by MATLAB's and Python's built-in solvers, whether or not the right-hand side depends explicitly on t . By following this format you can easily switch to using those built-in solvers, in case your own solver is not efficient enough.

3. Nonstiff and stiff problems

Theory. In the beginning of this instruction set, you used both the *explicit Euler method*,

$$y_{n+1} = y_n + hf(t_n, y_n)$$

and the *implicit Euler method*,

$$y_{n+1} = y_n + hf(t_{n+1}, y_{n+1}).$$

In an **implicit** method, every step is (far) more expensive due to the necessary (nonlinear) equation solving. This extra expense can pay off, however, **if one can take much longer steps** with the implicit method.

This happens in **stiff problems**. Because all explicit methods have bounded stability regions, the maximum stable step size is limited (see lecture notes). A well designed implicit method, however, can have an *unbounded stability region*, permitting much larger steps without losing accuracy. This makes up for the extra work incurred by equation solving. The explicit and implicit Euler methods are the prototypical examples of this.

The **van der Pol equation**,

$$\begin{aligned}y_1' &= y_2 \\y_2' &= \mu \cdot (1 - y_1^2) \cdot y_2 - y_1,\end{aligned}$$

models an electric oscillator circuit. The solution is periodic, with a period of approximately 2μ . The system may be stiff or nonstiff depending on the parameter μ . Use the initial condition $y(0) = [2 \ 0]^T$.

Task 3.1 Solve the van der Pol equation for $\mu = 100$ on the interval $[0, 2\mu]$ using your own explicit, adaptive RK34 code. Plot the solution component y_2 as function of time. Further, plot y_2 as a function of y_1 (the phase portrait). In the latter plot, try modifying the initial values, and check that the solution always tends to the same oscillation. Unlike the Lotka–Volterra equation, where you get different orbits, the van der Pol equation only has a single closed orbit, known as a **limit cycle**.

Task 3.2 We are now going to explore stiffness, and how it depends on μ . In all computations, use the initial condition $y(0) = [2 \ 0]^T$, and for every given μ , solve the problem on the time interval $[0, 0.7\mu]$, still using your own adaptive solver.

Solve the problem for the “E6 series” of values of μ , i.e., solve the problem for $\mu = 10, 15, 22, 33, 47, 68, 100, 150, 220, 330, 470, 680, 1000$. You may have to cut this series short if your code takes exceedingly long time to solve for say $\mu = 1000$. If the computation suddenly takes a suspiciously short time, check the results: “values” such as 10^{87} or NaN indicate that the adaptivity broke down and a much too large step size was used. Then we are outside the stability region, and both the RK approximations blow up in unpredictable ways. This can happen due to a too large initial step size, and a too stiff problem. Such results are not proper approximations.

Collect data by recording how many steps your solver needs to complete the integration in each case. Plot the total number of steps N as a function of μ in a loglog diagram. Use a suitable tolerance for all computations, at least $\text{TOL} = 10^{-6}$. The increase in the number of steps needed is proportional to the stiffness. How does stiffness depend on μ ?

Task 3.3 Read the documentation in MATLAB (using `help` and other sources) on how to use the stiff solver `ode15s` to solve differential equations. In Python, use the function `scipy.integrate.solve_IVP` and specify the parameter `method='BDF'`. This is roughly equivalent to `ode15s`. Repeat the experiments from Tasks 3.2 using one of these solvers, using the same data as before. What happens? Which code performs better, and why? When you plot the number of steps as a function of μ , how does the graph differ from the one you got with your own nonstiff RK34 code? Can you run the built-in solver for $\mu = 10,000$ and higher? (Don't even think of doing that with your own *explicit* code.) Does it take longer wall-clock time to solve a problem with μ large?

Please note that the Matlab/Python solvers do not outperform your adaptive RK34 method because they are “better implemented”, it is because they use implicit methods. We could do that as well, and the performance would be similar. But setting up an *implicit* Runge-Kutta method requires a bit more work, and so does getting a proper error estimate via an embedded scheme. If time permits, you can try applying your non-adaptive implicit Euler (or trapezoidal rule) method. It will likely work for any μ , but due to the lack of adaptiveness it will fail to pick up the parts where the solution changes very quickly, unless the step size is very small. (Note that this is not a question of stability but rather accuracy.)

Bonus task (do if interested): BDF. Implementing implicit Runge-Kutta methods with error estimates in a good way is difficult. It is even more difficult to properly implement multistep methods with adaptive step size control. Doing either of this is out of the scope of this course. But implementing multistep methods with a *fixed* step size is easy. It requires a somewhat different code structure, because you now have to keep the approximations at several previous time steps in memory. (If you previously kept all of them, that also works, but it requires excessive memory in large applications.)

For this task, implement one of the BDF methods. See the lecture notes/course book for the formulas, there are also nice tables with coefficients on e.g. Wikipedia. Do not choose BDF1, because this is equivalent to the implicit Euler method. Try it on the linear test equation to verify that it works. Then apply it to the Lotka-Volterra and van der Pol problems and compare it to your one-step methods. Because BDF is implicit, you need to solve an equation in each time step. For the linear test equation, you can use `scipy.linalg.solve` or MATLAB's backslash. For the other

nonlinear problems, you need to apply some iterative algebraic solver. We will talk more about this in the in connection to boundary value problems, but for now try e.g. the MATLAB function `fsolve` or the Python equivalent `scipy.optimize.fsolve`. (Check the documentation for syntax!)